

A photograph of a large herd of cows in a lush green field. In the foreground, a white cow stands on a grassy verge next to a paved road. A blue and yellow bicycle is leaning against a white post next to the cow. The background shows a rolling green hill with scattered trees and more cows grazing.

# **Real-Time Programming TI00AA55**

**Lecture 11 - 14.04.2015**

**Jarkko.Vuori@metropolia.fi**



# Real time signals 1

- The conventional signals are not queued and the signal handler cannot take parameters other than the signal number
- Current POSIX.1 incorporates also earlier POSIX.1b (real-time facilities) that defines real-time signals. They have additional features like:

1. Queueing of signals
2. Passing information to the signal handler

- Let's recall the definition of structure `sigaction`:

```
struct sigaction {  
    void (*sa_handler)(int signo);  
    sigset_t sa_mask;  
    int sa_flags;  
    → void (*sa_sigaction)  
        (int, siginfo_t *, void *); //rth  
}
```

- The last member is a so called real-time signal handler. It is an option to the first member. Real time signal handler is used, if the flag `SA_SIGINFO` is set in the member `sa_flags`. Then the prototype of the handler must be:  

```
void handler(int signo,  
             siginfo_t *siginfo,  
             void *context);
```
- Remark. The context parameter is not currently used
- A programmer can use real-time signals for his own purposes
- On the other hand, some new real-time APIs use them

# Real time signals 2

- Type `siginfo_t` (second parameter) is defined as follows:

```
typedef struct {
    int si_signo;
    int si_code;
    union sigval si_value
} siginfo_t;
```
- The member `si_code` of structure `siginfo_t` can have the following values:
  - `SI_USER`
  - `SI_QUEUE`
  - `SI_TIMER`
  - `SI_ASYNCIO`
  - `SI_MESGQ`
- Value `SI_USER` means that signal was sent with the conventional functions `kill`, `raise` or `abort`. In this case the member `si_value` has no meaning
- Value `SI_QUEUE` means that signal was sent with the new `sigqueue` function (see page 5)
- Value `SI_TIMER` means, that signal was sent because of the expiration of timer (so called real time timer, see page 14)
- Value `SI_ASYNCIO` means that signal was sent because of completion of asynchronous i/o
- Value `SI_MESGQ` means that message has arrived to the message queue

# Real time signals 3

- Real time signal handler can take application specific information in the member `si_value` of the parameter `siginfo`
  - The type of this member is union that makes it possible to pass `int` data or `void*` data
  - Actually the latter provides possibility to pass the pointer to what ever information to the signal handler
  - In this way we can avoid using of global variable when passing information between main process and signal handler
- The type union `sigval` (the type of member `si_value` of the structure `siginfo_t`) is defined as follows:

```
union sigval {  
    int    sival_int;  
    void *sival_ptr;  
}
```
- If signal was sent because of expiration of the timer, the parameter value that is received is determined when the timer was created with the function `timer_create` (see page 9)
- If the signal was sent because of completion of an asynchronous i/o, the parameter value contains the information that was set when i/o was initiated with the `aio_read` or `aio_write` function call
- If the signal was sent because a message has arrived to the message queue, the parameter value contains the information that was set with the function `mq_notify` (see lecture 13)

# Real time signals 4

- To send a real time signal from one process to another we need a new function
  - because it is not possible to pass parameter using function kill
  - Another reason is that the queueing system of signals does not work with kill function
- The new signal sending function is `sigqueue`
  - It guarantees that signals are queued
  - It must be used for real time signals
- Remember too that the signal queueing is working only if the signal handler for that signal is set with the function `sigaction` and with the flag `SA_SIGINFO` in the `sa_flags` member of `struct sigaction`
- Real time signals have their own signal numbers between constants `SIGRTMIN` – `SIGRTMAX`
- The prototype of the function `sigqueue` is

```
int sigqueue(pid_t pid, int signo, const union sigval value);
```
- Remark 1. If `_POSIX_REALTIME_SIGNALS` is defined, the implementation supports real time signals
- Remark 2. The programmer can send whatever value (int or void\*) to the real time signal handler
- Remark 3. It is dangerous to pass a pointer as a parameter to a signal handler from another process, because process cannot access the address space of another process

# Clocks and interval timers

- Clock means a counter, that is incremented at fixed time intervals
- Interval timer means a counter that kernel decrements at fixed intervals based on the clock and sends a signal when counter becomes 0, i.e. when a specified time has elapsed
- Unix systems have many different clocks and timers
- We have learned only most elementary clock that is requested by `time` function and most elementary timer that is requested by `alarm` function
  - This timer sends SIGALRM signal
- The disadvantage of this clock is a poor resolution (1s)
- The disadvantages of this timer is a poor resolution (1s) and the dependency on the SIGALRM
  - It does not restart automatically and you cannot have several timers at the same time
- Next we learn the most “advanced” and latest clocks and timers that POSIX.1-2008 recommends
- New timer API uses real-time signals

# Clocks 1

- Conventional functions `time`, `gettimeofday`, `clock` and `times` can be used to measure different times (wall clock time and processor time (see page 15))
- The latest **POSIX.1-2008 recommends** the following method. Clocks have a type `clockid_t`
  - A clock can be system wide or process wide clock
- There are the following clock ids and corresponding clocks in the system:
  - `CLOCK_REALTIME` ("wall clock time")
  - `CLOCK_MONOTONIC` (monotonic time)
  - `CLOCK_PROCESS_CPUTIME_ID` (proc cpu time)
  - `CLOCK_THREAD_CPUTIME_ID` (thread cpu time)
- `CLOCK_REALTIME` and `CLOCK_MONOTONIC` are system wide
  - Others are process or thread wide
- Everybody can read, but only super user can set these clocks
- The data type `struct timespec` represents times
- The time is even more accurate than in the `struct timeval`
  - The data member `tv_nsec` indicates nanoseconds
- The definition of `struct timespec` is:

```
struct timespec {
    time_t tv_sec;    // seconds
    long   tv_nsec;   // nanoseconds
};
```

Note that the `tv_nsec` is long in order to ensure that  $10^9$  fits the variable

  - because there are  $10^9$  nanoseconds in one second
  - Long is at least 32 bit in all implementations (range  $\approx -2 \cdot 10^9 \dots 2 \cdot 10^9$ )

# Clocks 2

- Functions that are used to handle these clocks are:

```
int clock_settime(clockid_t clock_id,  
const struct timespec *tp);  
clock_gettime(clockid_t clock_id,  
struct timespec *tp);  
clock_getres(clockid_t clock_id,  
struct timespec *tp);
```

- Return value is 0, if the function returns successfully
- Return value is -1, if the function execution has failed
- CLOCK\_REALTIME and CLOCK\_MONOTONIC measure time (in nanoseconds) since the Epoch 1.1.1970 at 00:00:00.
- The value that the function `clock_gettime` has produced in the member `tv_sec` of it's parameter can be used in a similar way than the value produced by the function `time`

- Example

```
struct timespec now;  
clock_gettime(CLOCK_REALTIME, &now);  
printf("%s", ctime(&now.tv_sec));
```

- CLOCK\_PROCESS\_CPUTIME\_ID and CLOCK\_THREAD\_CPUTIME\_ID measure the cpu time (in nanoseconds) used by the process or thread since the process or thread was started
- Remark 3. To be able to use these functions you need **to link real time library** (gcc needs `-lrt`)



# TMR Interval timers

- The latest POSIX.1-2008 recommends timers represented here (TMR timers)
- They provide a possibility to create several timers in the application
- Timers are based on the clocks (page 7)
- The user can choose, what clock the timer is based on
- The structure `struct itimerspec` is needed when TMR timers are used. The definition is as follows:

```
struct itimerspec {  
    struct timespec it_value; //timer expiration  
    struct timespec it_interval; //timer period  
};
```

- Functions you need to use these timers:  

```
int timer_create(clockid_t clock_id,  
                struct sigevent *evp,  
                timer_t *timerid);  
int timer_settime(timer_t timerid,  
                  int flags,  
                  const struct itimerspec *value,  
                  struct itimerspec *value);  
int timer_gettime(timer_t timerid,  
                  struct itimerspec *value);
```
- Because realtime signals are usually used with TMR timers lets recall the definition of the following structures:
  - `struct sigaction`
  - `siginfo_t`
  - `struct sigevent`
  - `union signal`

The definitions of these types are on the page 12

# Example of TMR interval timer

```
// Prototype of the signal handler
void rt_handler(int signo, siginfo_t *info, void* p ) {
    printf("Timer expired\n");
}

int main(void) {
    timer_t timerid;
    struct sigevent sevent;
    struct sigaction sa;
    struct itimerspec timerstruct;
    int i;

    // Install real time signal handler
    sigemptyset(&sa.sa_mask);
    sa.sa_flags = SA_SIGINFO;
    sa.sa_sigaction = rt_handler;
    sigaction(SIGRTMAX, &sa, NULL);
    // Construct a structure that specifies what signal is generated
    // and how signal handler is called
    sevent.sigev_notify = SIGEV_SIGNAL;
    sevent.sigev_signo = SIGRTMAX;
    sevent.sigev_value.sival_int = 2; // integer parameter
    // Create a timer instance
    timer_create(CLOCK_REALTIME, &sevent, &timerid);
    // Set the times for a timer
    timerstruct.it_value.tv_sec = 2;          // First time after 2 secs
    timerstruct.it_value.tv_nsec = 0;
    timerstruct.it_interval.tv_sec = 0;      // Do not repeat
    timerstruct.it_interval.tv_nsec = 0;     // Do not repeat
    timer_settime(timerid, 0, &timerstruct, NULL);
    // Do what ever
    for (i = 0 ; i < 6 ; i++) {
        sleep(1);
    }
    return 0;
}
```

# Absolute TMR interval timer

- The new counter value is loaded “automatically” by the kernel when the counter reaches the zero value
- However there is still an error that is called a drift
- To fully eliminate the drift TMR interval timer can use absolute time mode
- The second parameter (flags) of the function `timer_settime` has two options: 0 and `TIMER_ABSTIME`
- If the value of `TIMER_ABSTIME` is used, the structure members are interpreted as absolute time value
- See the example program in the web (`interval_timer_abs.c`) that demonstrates how absolute times are used to generate a signal in one second intervals without a drift

# Structure definitions

```
struct sigaction {  
    void (*sa_handler)(int signo);  
    sigset_t sa_mask;  
    int sa_flags;  
    void (*sa_sigaction)  
        (int, siginfo_t *, void *);  
};
```

```
typedef struct {  
    int si_signo;  
    int si_code;  
    union sigval si_value  
} siginfo_t;
```

```
struct sigevent {  
    int sigev_notify;           //notification type  
    int sigev_signo;           //signal number  
    union sigval sigev_value;  //signal value  
};
```

```
union sigval {  
    int sival_int;  
    void *sival_ptr;  
};
```

# More time functions and improvements

- We have learned earlier about different time functions:

```
struct tm *localtime(const time_t *timep);  
char *asctime(const struct tm *timep);  
char *ctime(const time_t *timep); // time zone  
struct tm *gmtime(const time_t *timep);
```

- The problem with these functions is that they are not reentrant
  - Function is called reentrant if it can be interrupted in the middle of its execution and then safely called again (“re-entered”) before its previous invocations complete execution
- To fix the problem POSIX:TSF extension defines corresponding thread safe functions:

```
struct tm *localtime_r(const time_t *timep, struct tm *result);  
char *asctime_r(const struct tm *timep, char *buff);  
char *ctime_r(const time_t *timep, char *buff); // time zone  
struct tm *gmtime_r(const time_t *timep, struct tm *result);
```

# gettimeofday

- There is also the function `gettimeofday` “between” time and `clock_gettime`. It uses `struct timeval`:

```
struct timeval {
    time_t tv_sec; // seconds since Epoch
    time_t tv_usec; // and microseconds
};
```

`<sys/times.h>`  
`int gettimeofday(struct timeval *tp, void *tzp);`  
(returns 0)
- Remark 1. This type is same that was used as a timeout in the function `select` (the last parameter)
  - In that context the time was used to express time interval (not the absolute time)

- Remark 2. The value in the member `tv_sec` after calling `gettimeofday` can be used in similar way than the value produced by the function `time`
- Example

```
struct timeval now;
gettimeofday(&now, NULL);
printf("%s", ctime(&now.tv_sec));
```
- Remark. The new version of POSIX.1-2008 recommends rather `clock_gettime` (see page 7)



# Conventional clocks 1

- ANSI C defines a function `clock`  
`<time.h>`  
`clock_t clock(void);`
- The type of the reading of this conventional clock is `clock_t`
- This function returns the **cpu time** of the process that is used so far
  - POSIX requires that time unit should be a micro second (`CLOCKS_PER_SEC` is 1000000)
- UNIX specification defines another function `times`, that can be used to measure different process times: total time (wall clock time), user cpu time and system cpu time
  - The sum of the last two times is the total cpu time
- The prototype of the function is:  
`<times.h>`  
`clock_t times(struct tms *buff);`
- Return value is the total time, and the structure `struct tms` is defined as follows:

```
struct tms {  
    clock_t tms_utime; //user cpu time used so far  
    clock_t tms_stime; //system cpu time used  
    clock_t tms_cutime; //user cpu time of children  
    clock_t tms_cstime; //system time of children  
}
```
- The function returns the value of tick counter since the booting
  - Ticks per second can be determined with the function `sysconf(_SC_CLK_TCK)`

# Conventional clocks 2

- Remark. The two functions on the previous page `clock` and `times` both return the time of type `clock_t`
  - However, these times mean different things and even time units are different
  - The function `clock` returns the processor time of the process and the function `times` return the wall clock time
  - The time units are explained on the right
- ISO C

```
clock_t clock(void);
```

CPU-time (microseconds)  
The unit is specified with the constant `CLOCKS_PER_SEC` (is 1000000)
- UNIX

```
clock_t times(struct tms *buff);
```

Wall clock time (1 / 100 seconds or ticks)
  - The unit is specified with the value returned by the function `sysconf _SC_CLK_TCK` as parameter (is 100)

# Delays (sleeping)

- Different sleep functions:
  - A. `unsigned sleep(unsigned seconds);`
    - Problems:
      1. resolution only one second
      2. uses SIGALRM signal
  - B. `int nanosleep(const struct timespec *requested, struct timespec *remaining);`
    - This function has no interaction with the SIGALRM
    - Resolution is nanosecond
- Both functions return, if time has elapsed or signal is delivered and signal handler returns
- Example (Delay of 0.5 seconds)

```
struct timespec delay =  
{ 0, 500000000 };  
nanosleep(&delay, NULL);
```
- Remark 1. If signals are delivered to the process and you want to make sure that the required delay really happens, you need to write a loop using the principle in much the same way we did with the sleep function
- Remark 2. There is also function `usleep`, that takes the delay as microseconds
  - The delay must not be a second or longer
  - This means that parameter needs to be less than 1 000 000
  - For example `usleep(500000)` causes 0.5 second delay
- The use of this function is not recommended

# POSIX:XSI Interval timers 1

- POSIX:XSI also defines interval timers
- There are three different kind of interval timers:
  - ITIMER\_REAL (based on the "wall" clock)
  - ITIMER\_PROF (based on the cpu time)
  - ITIMER\_VIRTUAL (based on the system time)
- The time intervals are expressed with the structures itimerval:
- Note that the time is defined in the accuracy of microseconds because `struct timeval` is used
- Function `getitimer` can be used to read the current value of interval timer and `setitimer` can be used to start and stop the interval timer

```
struct itimerval {  
    struct timeval it_value; // time until  
    next expiration  
    struct timeval it_interval; // value to  
    reload into the timer;  
};
```

```
int getitimer(int which, struct  
itimerval *value);  
int setitimer(int which, const struct  
itimerval *value, struct itimerval  
*ovalue);
```

# POSIX:XSI interval timers 2

- When timer is set using the function `setitimer`, it expires when the time that was set in the member `it_value` of parameter `value` has elapsed
  - Kernel decrements this parameter member according the rule given as the first parameter (which)
  - When time in the member `it_value` becomes zero, kernel sends a signal `SIGALRM` to the process
- Kernel loads automatically the member `it_value` with the member `it_interval`, so that timer starts from the begin without any interaction from the user program
- The process can query the current value of the counter (member `it_value`) at any time with the function `getitimer`
- Remark. The new version of POSIX.1-2008 **recommends rather TMR interval timer** (see page 9)

# Clocks (summary)

- POSIX:TMR
- CLOCK\_REALTIME
  - Measures wall clock time
- CLOCK\_MONOTONIC
  - Measures wall clock time using monotonic clock
  - (Time changes have no effect)
- CLOCK\_PROCESS\_CPUTIME\_ID
  - Measures total cpu-time of the process
- CLOCK\_THREAD\_CPUTIME\_ID
  - Measures total cpu-time of the thread
- Use structure `struct timespec` (seconds and nanoseconds)
- Functions: `clock_settime`, `clock_gettime` and `clock_getres`



# Timers (summary)

- POSIX:TMR

- It is possible to create several timers
- You can choose the clock which the timer is based on
- The structure `struct itimerspec` is used that contains two members each of which is of type `struct timespec` (seconds and nanoseconds)
- Functions: `timer_create`, `timer_settime` and `timer_gettime`
- Function `timer_create` produces timer id
- User can specify the signal (real time signal)

- POSIX:XSI

- ITIMER\_REAL
  - Is based on the wall clock
- ITIMER\_PROF
  - Is based on the cpu-time
- ITIMER\_VIRTUAL
  - Is based on the system time
- Structure `struct itimerval` is used, that contains two members, each of which is of type `struct timeval` (seconds and microseconds)
- Functions: `getitimer` and `setitimer`
- Signal SIGALRM is generated
- Conventional: Function `alarm` and signal SIGALRM

# Conventional time functions (summary)

- `time`
  - Resolution is second
- `gettimeofday`
  - Resolution is micro second
  - Takes structure `struct timeval` as a parameter
- `clock`
  - Returns cpu-time (unit is microsecond, `CLOCKS_PER_SEC`)
- `times`
  - Produces user cpu time and system cpu time in the structure `struct tms`
  - Returns wall clock time (unit is tick, `sysconf(_SC_CLK_TCK)`)

```
struct tms {  
    clock_t tms_utime;  
    clock_t tms_stime;  
    clock_t tms_cutime;  
    clock_t tms_ctime;  
};
```